

ICCS413 - Lecture 18

# Ensemble Learning and Face Detection


*Sunsern Cheamanunkul*



Mahidol University  
International College




# Face Detection



# Basic Concepts

- Slide a window across the image.
- Test each window for a face



# Challenges

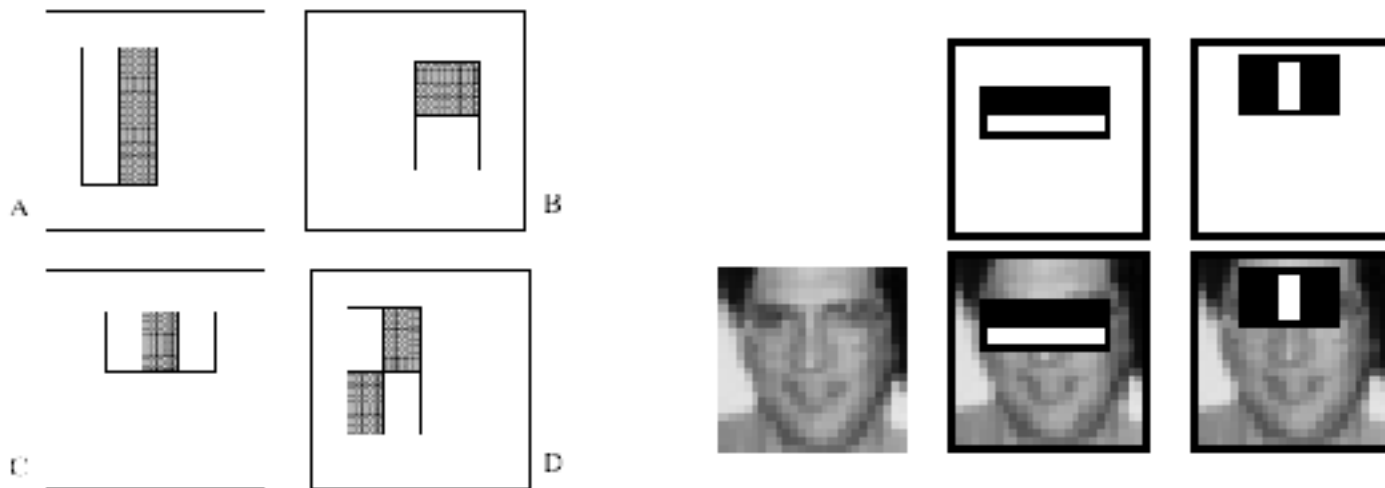
- Sliding window detector must evaluate  $\geq 10k$  location/scale combinations
- Faces are rare: 0-10 per image
  - For computational efficiency, we should try to spend as little time as possible on the non-face windows
  - A hi-res image has  $\sim 10^6$  pixels and a comparable number of candidate face locations
  - To avoid having false positives in every image, the false positive rate has to be very small (less than  $10^{-6}$ )

# Viola-Jones Face Detector

- **Robust** - High true positive rate while maintaining low false positive rate
- **Real-time** - For practical applications at least 2 frames per second must be processed.
- **Face Detection** - not recognition. The goal is to distinguish faces from non-faces (face detection is the first step in the identification process)


# Features

- Four basic types
  - Simple to calculate
  - The white areas are subtracted from the black ones.
  - A special representation of the sample called the integral image makes feature extraction faster.



# Integral Images

- Summed area tables



- A representation that means any rectangle's values can be calculated in four accesses of the integral image.

# Fast Pixel Summation





Figure 3: The sum of the pixels within rectangle  $D$  can be computed with four array references. The value of the integral image at location 1 is the sum of the pixels in rectangle  $A$ . The value at location 2 is  $A + B$ , at location 3 is  $A + C$ , and at location 4 is  $A + B + C + D$ . The sum within  $D$  can be computed as  $4 + 1 - (2 + 3)$ .



# Feature extraction

- Features are extracted from each sub-window of the image.
  - The base size for a sub-window is 24 by 24 pixels.
  - Each of the four feature types are scaled and shifted across all possible combinations
- In a 24 pixel by 24 pixel sub window there are ~160,000 possible features to be calculated.




# Challenges of the Learning Task

- High number of features (~160,000)
- Only a few hundreds training examples
- Which features to use?
- How to avoid overfitting?
- Boosting!
  - Combine many weak classifiers into a strong classifier

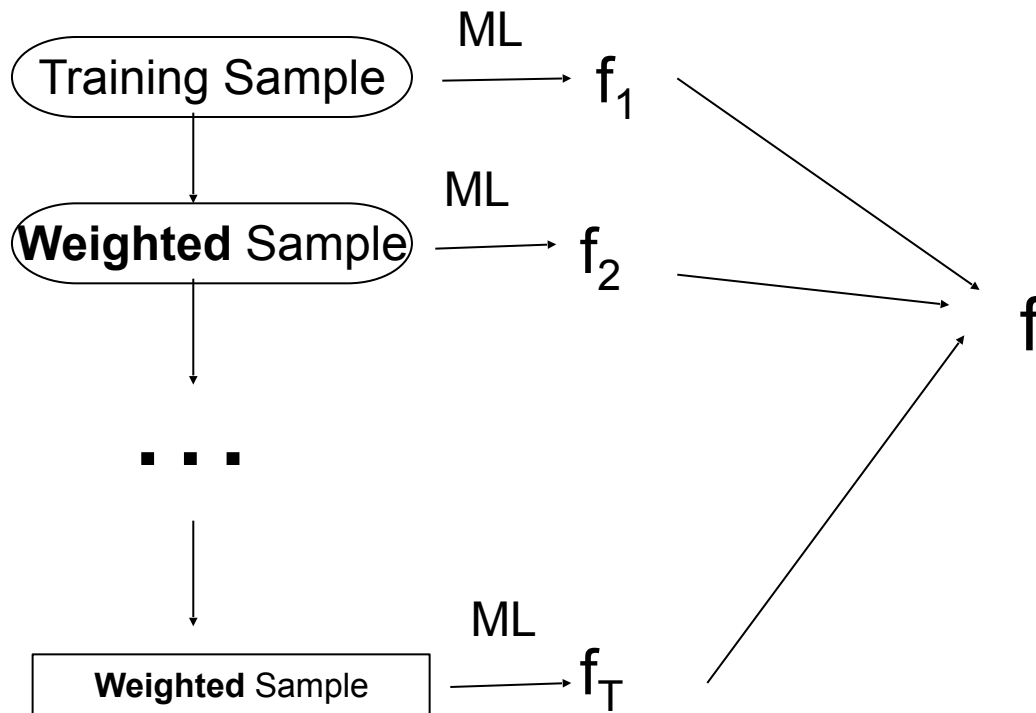
# Ensemble Learning

- a process by which multiple models, such as classifiers or experts, are strategically generated and combined to solve a particular computational intelligence problem such as classification, prediction, function approximation.
- Examples:
  - Bagging
  - Boosting

# Bagging



# Boosting



# Overview of boosting

- Introduced by Schapire and Freund in 1990s.
- “Boosting”: convert a weak learning algorithm into a strong one.
- Main idea: Combine many “weak” classifiers to produce a “strong” classifier.
- Algorithms:
  - **AdaBoost**: adaptive boosting
  - Gentle Boost
  - BrownBoost
  - Logit Boost
  - etc.

# Intuition

- Train a set of weak hypotheses:  $h_1, \dots, h_T$ .
- The combined hypothesis  $H$  is a **weighted** majority vote of the  $T$  weak hypotheses.
  - Each hypothesis  $h_t$  has a weight  $\alpha_t$ .

$$H(x) \equiv \text{sgn} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$$

- During the training, focus on the examples that are misclassified.
  - At round  $t$ , example  $x_i$  has the weight  $D_t(i)$ .

# Basic Setting

- Binary classification problem
- Training data:

$(x_1, y_1), \dots, (x_m, y_m)$ , where  $x_i \in X, y_i \in Y = \{-1, 1\}$

- $D_t(i)$ : the weight of  $x_i$  at round  $t$ .  $D_1(i) = 1/m$ .
- A learner  $L$  that finds a weak hypothesis  $h_t: X \rightarrow Y$  given the training set and  $D_t$
- The error of a weak hypothesis  $h_t$ :

$$\epsilon_t \equiv \Pr_{i \sim D_t} [h_t(x_i) \neq y_i] \equiv \sum_{i: h_t(x_i) \neq y_i} D_t(i)$$



# AdaBoost algorithm

Given:  $(x_1, y_1), \dots, (x_m, y_m)$  where  $x_i \in X$ ,  $y_i \in Y = \{-1, +1\}$

Initialize  $D_1(i) = 1/m$ .

For  $t = 1, \dots, T$ :

- Train weak learner using distribution  $D_t$ .
- Get weak hypothesis  $h_t : X \rightarrow \{-1, +1\}$  with error

$$\epsilon_t = \Pr_{i \sim D_t} [h_t(x_i) \neq y_i].$$

- Choose  $\alpha_t = \frac{1}{2} \ln \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$ .
- Update:

$$\begin{aligned} D_{t+1}(i) &= \frac{D_t(i)}{Z_t} \times \begin{cases} e^{-\alpha_t} & \text{if } h_t(x_i) = y_i \\ e^{\alpha_t} & \text{if } h_t(x_i) \neq y_i \end{cases} \\ &= \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t} \end{aligned}$$

where  $Z_t$  is a normalization factor (chosen so that  $D_{t+1}$  will be a distribution).

Output the final hypothesis:

$$H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right).$$

# Strengths of AdaBoost

- It has no parameters to tune (except for the number of rounds)
- It is fast, simple and easy to program.
- It comes with a set of theoretical guarantee (e.g., training error, test error)
- Instead of trying to design a learning algorithm that is accurate over the entire space, we can focus on finding base learning algorithms that only need to be better than random.
- It can identify outliers: i.e. examples that are either mislabeled or that are inherently ambiguous and hard to categorize.

# Weakness of AdaBoost

- The actual performance of boosting depends on the data and the base learner.
- AdaBoost is sensitive to noise.
- When the number of outliers is very large, the emphasis placed on the hard examples can hurt the performance.  
→ “Gentle AdaBoost”, “BrownBoost”

# Bagging vs. Boosting


## (Freund and Schapire 1996)

- Bagging always uses resampling rather than reweighting.
- Bagging does not modify the distribution over examples or mislabels, but instead always uses the uniform distribution
- In forming the final hypothesis, bagging gives equal weight to each of the weak hypotheses


# Weak Classifier

- Decision stumps = decision tree with only a single root node
  - Certainly a very weak learner!
  - Say the attributes are real-valued
  - Decision stump algorithm looks at all possible thresholds for each attribute
  - Selects the one with the max information gain
  - Resulting classifier is a simple threshold on a single feature
    - Outputs a +1 if the attribute is above a certain threshold
    - Outputs a -1 if the attribute is below the threshold


# Boosting Example



# After 1st Rule




# After 2nd Rule






# After 3rd Rule



# Final Classifier

$$f_{\text{final}} = \text{sign} \left( 0.42 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} + 0.65 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} + 0.92 \begin{array}{|c|} \hline \text{blue} \\ \hline \end{array} \right)$$


The diagram shows three square regions, each divided by a vertical line. The first square has a blue left half and a pink right half. The second square has a blue left half and a pink right half. The third square has a blue top half and a pink bottom half. Each square is enclosed in green parentheses.




# Boosting with Decision Stumps

- Viola-Jones algorithm
  - With  $K$  attributes (e.g.,  $K = 160,000$ ) we have 160,000 different decision stumps to choose from
  - At each stage of boosting
    - given reweighted data from previous stage
    - Train all  $K$  (160,000) single-feature stumps
    - Select the single best classifier at this stage
    - Combine it with the other previously selected classifiers
    - Reweight the data
    - Learn all  $K$  classifiers again, select the best, combine, reweight
    - Repeat until you have  $T$  classifiers selected
  - Very computationally intensive
    - Learning  $K$  decision stumps  $T$  times
    - E.g.,  $K = 160,000$  and  $T = 1000$

# Training error as a function of boosting iterations



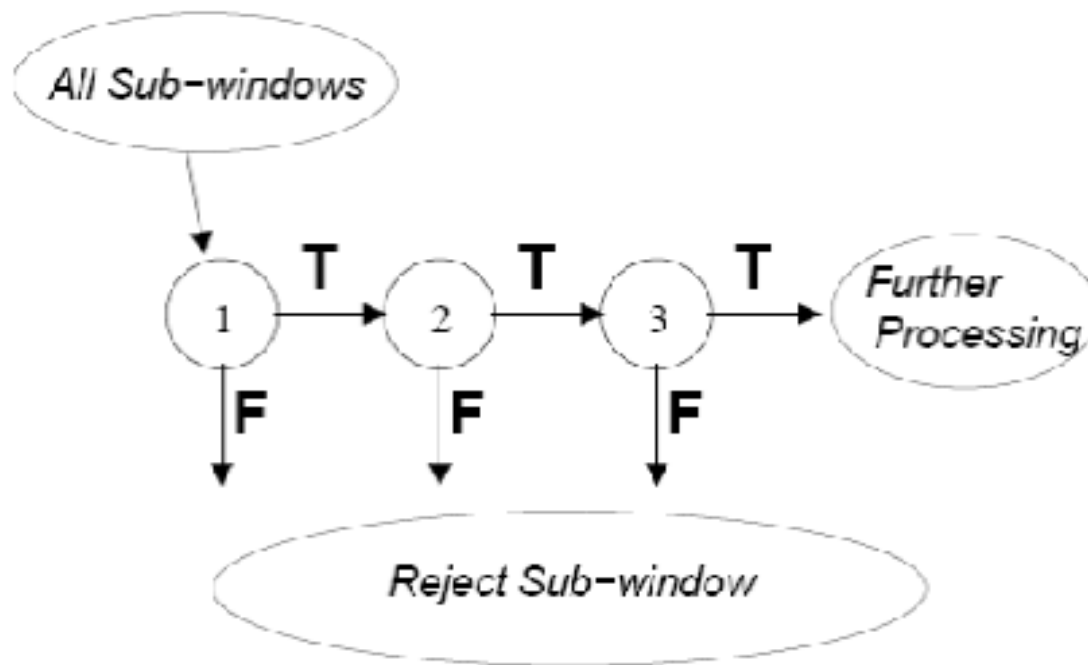
# Top features identified by Boosting



# Real-time Face Detection

- Basic classifier operates on 24 x 24 subwindows
- Scaling:
  - Scale the detector (rather than the images)
  - Features can easily be evaluated at any scale
  - Scale by factors of 1.25
- Location:
  - Move detector around the image (e.g., 1 pixel increments)
- Final Detections
  - A real face may result in multiple nearby detections
  - Post-process detected subwindows to combine overlapping detections into a single detection

# Cascading




# Training Images

- Examples of 24x24 images with faces





# Face vs Nonface



# Results

